

In [3]: **import** pandas **as** pd

```
# Loading the dataset
file_path = "GlobalWeatherRepository.csv"
df = pd.read_csv(file_path)
```

In [5]: *# Printing dataset information*

```
print("\n Displaying basic information about the dataset:\n")
df.info()

# Printing first few rows
print("\n Displaying the first 5 rows of the dataset:\n")
print(df.head())

# Printing dataset size
print("\n The dataset contains", df.shape[0], "rows and", df.shape[1], "columns")

# Printing column names
print("\n Column names in the dataset:\n", df.columns.tolist(), "\n")

# Printing data types of columns
print("\n Data types of each column:\n")
print(df.dtypes)

# Checking missing values
print("\n Checking for missing values in the dataset:\n")
print(df.isnull().sum())

# Checking for duplicate rows
print("\n Number of duplicate rows in the dataset:", df.duplicated().sum(),)
```

Displaying basic information about the dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 56906 entries, 0 to 56905

Data columns (total 41 columns):

#	Column	Non-Null Count		Dtype
0	country	56906	non-null	object
1	location_name	56906	non-null	object
2	latitude	56906	non-null	float64
3	longitude	56906	non-null	float64
4	timezone	56906	non-null	object
5	last_updated_epoch	56906	non-null	int64
6	last_updated	56906	non-null	object
7	temperature_celsius	56906	non-null	float64
8	temperature_fahrenheit	56906	non-null	float64
9	condition_text	56906	non-null	object
10	wind_mph	56906	non-null	float64
11	wind_kph	56906	non-null	float64
12	wind_degree	56906	non-null	int64
13	wind_direction	56906	non-null	object
14	pressure_mb	56906	non-null	float64
15	pressure_in	56906	non-null	float64
16	precip_mm	56906	non-null	float64
17	precip_in	56906	non-null	float64
18	humidity	56906	non-null	int64
19	cloud	56906	non-null	int64
20	feels_like_celsius	56906	non-null	float64
21	feels_like_fahrenheit	56906	non-null	float64
22	visibility_km	56906	non-null	float64
23	visibility_miles	56906	non-null	float64
24	uv_index	56906	non-null	float64
25	gust_mph	56906	non-null	float64
26	gust_kph	56906	non-null	float64
27	air_quality_Carbon_Monoxide	56906	non-null	float64
28	air_quality_Ozone	56906	non-null	float64
29	air_quality_Nitrogen_dioxide	56906	non-null	float64
30	air_quality_Sulphur_dioxide	56906	non-null	float64
31	air_quality_PM2.5	56906	non-null	float64
32	air_quality_PM10	56906	non-null	float64
33	air_quality_us-epa-index	56906	non-null	int64
34	air_quality_gb-defra-index	56906	non-null	int64
35	sunrise	56906	non-null	object
36	sunset	56906	non-null	object
37	moonrise	56906	non-null	object
38	moonset	56906	non-null	object
39	moon_phase	56906	non-null	object
40	moon_illumination	56906	non-null	int64

dtypes: float64(23), int64(7), object(11)

memory usage: 17.8+ MB

Displaying the first 5 rows of the dataset:

	country	location_name	latitude	longitude	timezone	\
0	Afghanistan	Kabul	34.52	69.18	Asia/Kabul	
1	Albania	Tirana	41.33	19.82	Europe/Tirane	

2	Algeria	Algiers	36.76	3.05	Africa/Algiers
3	Andorra	Andorra La Vella	42.50	1.52	Europe/Andorra
4	Angola	Luanda	-8.84	13.23	Africa/Luanda

	last_updated_epoch	last_updated	temperature_celsius	\
0	1715849100	2024-05-16 13:15	26.6	
1	1715849100	2024-05-16 10:45	19.0	
2	1715849100	2024-05-16 09:45	23.0	
3	1715849100	2024-05-16 10:45	6.3	
4	1715849100	2024-05-16 09:45	26.0	

	temperature_fahrenheit	condition_text	...	air_quality_PM2.5	\
0	79.8	Partly Cloudy	...	8.4	
1	66.2	Partly cloudy	...	1.1	
2	73.4	Sunny	...	10.4	
3	43.3	Light drizzle	...	0.7	
4	78.8	Partly cloudy	...	183.4	

	air_quality_PM10	air_quality_us-epa-index	air_quality_gb-defra-index	\
0	26.6	1	1	
1	2.0	1	1	
2	18.4	1	1	
3	0.9	1	1	
4	262.3	5	10	

	sunrise	sunset	moonrise	moonset	moon_phase	moon_illumination
0	04:50 AM	06:50 PM	12:12 PM	01:11 AM	Waxing Gibbous	55
1	05:21 AM	07:54 PM	12:58 PM	02:14 AM	Waxing Gibbous	55
2	05:40 AM	07:50 PM	01:15 PM	02:14 AM	Waxing Gibbous	55
3	06:31 AM	09:11 PM	02:12 PM	03:31 AM	Waxing Gibbous	55
4	06:12 AM	05:55 PM	01:17 PM	12:38 AM	Waxing Gibbous	55

[5 rows x 41 columns]

The dataset contains 56906 rows and 41 columns.

Column names in the dataset:

```
['country', 'location_name', 'latitude', 'longitude', 'timezone', 'last_updated_epoch', 'last_updated', 'temperature_celsius', 'temperature_fahrenheit', 'condition_text', 'wind_mph', 'wind_kph', 'wind_degree', 'wind_direction', 'pressure_mb', 'pressure_in', 'precip_mm', 'precip_in', 'humidity', 'cloud', 'feels_like_celsius', 'feels_like_fahrenheit', 'visibility_km', 'visibility_miles', 'uv_index', 'gust_mph', 'gust_kph', 'air_quality_Carbon_Monoxide', 'air_quality_Ozone', 'air_quality_Nitrogen_dioxide', 'air_quality_Sulphur_dioxide', 'air_quality_PM2.5', 'air_quality_PM10', 'air_quality_us-epa-index', 'air_quality_gb-defra-index', 'sunrise', 'sunset', 'moonrise', 'moonset', 'moon_phase', 'moon_illumination']
```

Data types of each column:

country	object
location_name	object
latitude	float64
longitude	float64

timezone	object
last_updated_epoch	int64
last_updated	object
temperature_celsius	float64
temperature_fahrenheit	float64
condition_text	object
wind_mph	float64
wind_kph	float64
wind_degree	int64
wind_direction	object
pressure_mb	float64
pressure_in	float64
precip_mm	float64
precip_in	float64
humidity	int64
cloud	int64
feels_like_celsius	float64
feels_like_fahrenheit	float64
visibility_km	float64
visibility_miles	float64
uv_index	float64
gust_mph	float64
gust_kph	float64
air_quality_Carbon_Monoxide	float64
air_quality_Ozone	float64
air_quality_Nitrogen_dioxide	float64
air_quality_Sulphur_dioxide	float64
air_quality_PM2.5	float64
air_quality_PM10	float64
air_quality_us-epa-index	int64
air_quality_gb-defra-index	int64
sunrise	object
sunset	object
moonrise	object
moonset	object
moon_phase	object
moon_illumination	int64
dtype:	object

Checking for missing values in the dataset:

country	0
location_name	0
latitude	0
longitude	0
timezone	0
last_updated_epoch	0
last_updated	0
temperature_celsius	0
temperature_fahrenheit	0
condition_text	0
wind_mph	0
wind_kph	0
wind_degree	0
wind_direction	0
pressure_mb	0

```

pressure_in      0
precip_mm        0
precip_in        0
humidity         0
cloud            0
feels_like_celsius 0
feels_like_fahrenheit 0
visibility_km     0
visibility_miles  0
uv_index         0
gust_mph         0
gust_kph         0
air_quality_Carbon_Monoxide 0
air_quality_Ozone 0
air_quality_Nitrogen_dioxide 0
air_quality_Sulphur_dioxide 0
air_quality_PM2.5 0
air_quality_PM10 0
air_quality_us-epa-index 0
air_quality_gb-defra-index 0
sunrise          0
sunset           0
moonrise         0
moonset          0
moon_phase       0
moon_illumination 0
dtype: int64

```

Number of duplicate rows in the dataset: 0

Analysis from the above output

The dataset consists of **56,906 rows** and **41 columns**, containing **daily weather data** for various global locations. It includes a variety of meteorological and environmental parameters:

- **Temperature, Wind Speed, Pressure, Precipitation, Humidity, and Air Quality**
- **Astronomical Data** such as **sunrise, sunset, and moon phase**

Data Structure

- The dataset is **well-structured** with appropriate **data types**:
 - **Numerical features**: Stored as float or int64
 - **Categorical attributes**: Stored as object (e.g., weather conditions, wind direction)
 - **Timestamp column (last_updated)**: Currently in object format and needs conversion for **time series analysis**
- **No missing values** or **duplicate rows**, meaning **minimal preprocessing** is required.

Data Cleaning & Preprocessing

- **Redundant Columns:** Some features are duplicated in different units:
 - **Temperature** (Celsius & Fahrenheit)
 - **Wind Speed** (mph & kph)
 - These will be **dropped to avoid duplication**.
- **Categorical Feature Consistency:**
 - Weather condition descriptions and **wind direction** need to be checked for **uniformity**.

```
In [8]: df['last_updated'] = pd.to_datetime(df['last_updated'])
print("Converted 'last_updated' to datetime format.")
```

Converted 'last_updated' to datetime format.

Data Cleaning: Removing Redundant Columns

To eliminate **duplicate information** and ensure **consistency** in the dataset, we removed **redundant columns** that provided the same data in different units. The following columns were dropped:

- **Temperature:** temperature_fahrenheit (retained temperature_celsius)
- **Pressure:** pressure_in (retained pressure_mb)
- **Precipitation:** precip_in (retained precip_mm)
- **Feels Like Temperature:** feels_like_fahrenheit (retained feels_like_celsius)
- **Visibility:** visibility_miles (retained visibility_km)
- **Wind Gust:** gust_mph (retained gust_kph)

Reasons for Removing These Columns

- **Ensured uniformity** by keeping only **metric units**.
- **Reduced memory usage**, making the dataset more efficient.
- **Prevented confusion** during analysis by maintaining **a single unit system**.

This step optimizes the dataset, making it **cleaner** and **more efficient** for **exploratory data analysis (EDA)** and **forecasting models**.

```
In [11]: # List of columns to drop (only the ones that exist)
columns_to_drop = ['temperature_fahrenheit', 'pressure_in', 'precip_in',
                  'feels_like_fahrenheit', 'visibility_miles', 'gust_mph']

# Drop only existing columns to avoid KeyError
df.drop(columns=columns_to_drop, inplace=True)

print(f"Dropped redundant columns: {columns_to_drop}\n")
```

Dropped redundant columns: ['temperature_fahrenheit', 'pressure_in', 'precip_in', 'feels_like_fahrenheit', 'visibility_miles', 'gust_mph']

```
In [13]: # Sort dataset by datetime column
df = df.sort_values(by='last_updated')
print("Sorted dataset by 'last_updated'.")
```

Sorted dataset by 'last_updated'.

```
In [15]: # Check unique values in categorical columns
print("Unique values in 'condition_text':\n", df['condition_text'].unique())
print("Unique values in 'wind_direction':\n", df['wind_direction'].unique())
print("Unique values in 'moon_phase':\n", df['moon_phase'].unique())
```

```
Unique values in 'condition_text':
['Clear' 'Fog' 'Overcast' 'Moderate or heavy rain with thunder'
 'Patchy rain nearby' 'Mist' 'Partly cloudy' 'Partly Cloudy' 'Sunny'
 'Moderate or heavy rain shower' 'Light rain' 'Moderate rain'
 'Light drizzle' 'Thundery outbreaks in nearby'
 'Patchy light rain in area with thunder' 'Patchy light rain with thunder'
 'Moderate rain at times' 'Light rain shower' 'Cloudy'
 'Heavy rain at times' 'Patchy light rain' 'Patchy light drizzle'
 'Thundery outbreaks possible' 'Patchy rain possible'
 'Moderate or heavy rain in area with thunder' 'Heavy rain'
 'Torrential rain shower' 'Freezing fog' 'Moderate or heavy snow showers'
 'Light sleet' 'Blizzard' 'Moderate snow' 'Light snow'
 'Light sleet showers' 'Light freezing rain' 'Heavy snow' 'Blowing snow'
 'Patchy heavy snow' 'Light snow showers' 'Moderate or heavy sleet'
 'Patchy light snow' 'Patchy moderate snow' 'Freezing drizzle'
 'Moderate or heavy snow in area with thunder' 'Patchy snow nearby'
 'Patchy snow possible' 'Patchy light snow in area with thunder']
Unique values in 'wind_direction':
['SW' 'N' 'E' 'S' 'ESE' 'SSW' 'WSW' 'SE' 'ENE' 'SSE' 'NE' 'NNE' 'NNW'
 'WNW' 'W' 'NW']
Unique values in 'moon_phase':
['Waxing Gibbous' 'Full Moon' 'Waning Gibbous' 'Last Quarter'
 'Waning Crescent' 'New Moon' 'Waxing Crescent' 'First Quarter']
```

After Analyzing the Above Output

After reviewing the **unique values** in the categorical data, we identified minor inconsistencies in the `condition_text` (weather conditions) column. Some variations, such as **"Partly cloudy" vs. "Partly Cloudy"** and **"Patchy rain nearby" vs. "Patchy rain possible"**, indicate slight formatting differences or similar meanings. To ensure uniformity, we will **convert all weather condition values to lowercase**. The `wind_direction` column was found to be **consistent**, with all values using standardized **uppercase abbreviations** (e.g., N, NE, NW). Similarly, the `moon_phase` column had **no formatting issues**, with all values correctly structured. These steps help **maintain data integrity** and ensure **cleaner categorical data** for further analysis.

```
In [18]: # Standardize weather condition text to lowercase
df['condition_text'] = df['condition_text'].str.lower()
print("Standardized 'condition_text' to lowercase for consistency.")
```

Standardized 'condition_text' to lowercase for consistency.

```
In [20]: print(df['condition_text'].unique())
```

['clear' 'fog' 'overcast' 'moderate or heavy rain with thunder'
'patchy rain nearby' 'mist' 'partly cloudy' 'sunny'
'moderate or heavy rain shower' 'light rain' 'moderate rain'
'light drizzle' 'thundery outbreaks in nearby'
'patchy light rain in area with thunder' 'patchy light rain with thunder'
'moderate rain at times' 'light rain shower' 'cloudy'
'heavy rain at times' 'patchy light rain' 'patchy light drizzle'
'thundery outbreaks possible' 'patchy rain possible'
'moderate or heavy rain in area with thunder' 'heavy rain'
'torrential rain shower' 'freezing fog' 'moderate or heavy snow showers'
'light sleet' 'blizzard' 'moderate snow' 'light snow'
'light sleet showers' 'light freezing rain' 'heavy snow' 'blowing snow'
'patchy heavy snow' 'light snow showers' 'moderate or heavy sleet'
'patchy light snow' 'patchy moderate snow' 'freezing drizzle'
'moderate or heavy snow in area with thunder' 'patchy snow nearby'
'patchy snow possible' 'patchy light snow in area with thunder']

Explanation of Data Cleaning & Preprocessing

We performed several **data cleaning steps** to prepare the dataset for analysis:

- **Converted** last_updated to **datetime** to enable **time-series analysis**.
- **Dropped redundant columns** that had duplicate information (e.g., **temperature in both Celsius & Fahrenheit**).
- **Sorted the dataset** by last_updated to ensure **chronological order**.
- **Standardized** condition_text to **lowercase** to eliminate **inconsistencies** in weather condition labels.

These steps ensure that the dataset is **clean, consistent, and optimized for analysis**.

Now, we are ready to move to the **Exploratory Data Analysis (EDA)** phase!

EDA

```
In [24]: print(df.describe())
```


	latitude	longitude	last_updated_epoch \
count	56906.000000	56906.000000	5.690600e+04
mean	19.136988	22.187380	1.728530e+09
min	-41.300000	-175.200000	1.715849e+09
25%	3.750000	-6.836100	1.722255e+09
50%	17.250000	23.320000	1.728554e+09
75%	40.400000	50.580000	1.734862e+09
max	64.150000	179.220000	1.741169e+09
std	24.477303	65.808904	7.355847e+06

	last_updated	temperature_celsius	wind_mph \
count	56906	56906.000000	56906.000000
mean	2024-10-10 05:31:21.892946176	22.278399	8.289013
min	2024-05-16 01:45:00	-24.900000	2.200000
25%	2024-07-29 15:15:00	17.100000	4.000000
50%	2024-10-10 12:45:00	25.100000	6.900000
75%	2024-12-22 15:45:00	28.500000	11.600000
max	2025-03-05 23:00:00	49.200000	1841.200000
std	NaN	9.647370	9.398604

	wind_kph	wind_degree	pressure_mb	precip_mm	...	\
count	56906.000000	56906.000000	56906.000000	56906.000000
mean	13.343804	169.668857	1014.164341	0.140864
min	3.600000	1.000000	947.000000	0.000000
25%	6.500000	80.000000	1010.000000	0.000000
50%	11.200000	160.000000	1013.000000	0.000000
75%	18.700000	258.000000	1018.000000	0.030000
max	2963.200000	360.000000	3006.000000	42.240000
std	15.123937	103.767926	13.831804	0.605114

	gust_kph	air_quality_Carbon_Monoxide	air_quality_Ozone	\
count	56906.000000	56906.000000	56906.000000	...
mean	19.176788	525.75107	63.465594	...
min	3.600000	-9999.000000	0.000000	...
25%	10.800000	223.600000	38.600000	...
50%	16.600000	323.750000	60.100000	...
75%	25.600000	500.700000	83.000000	...
max	2970.400000	38879.39800	480.700000	...
std	16.949292	954.16238	36.638077	...

	air_quality_Nitrogen_dioxide	air_quality_Sulphur_dioxide	\
count	56906.000000	56906.000000	...
mean	14.831329	11.301580	...
min	0.000000	-9999.000000	...
25%	0.925000	0.740000	...
50%	3.300000	2.220000	...
75%	15.910000	8.695000	...
max	427.700000	521.330000	...
std	26.244844	49.628464	...

	air_quality_PM2.5	air_quality_PM10	air_quality_us-epa-index	\
count	56906.000000	56906.000000	56906.000000	...
mean	25.130634	50.476481	1.711015	...
min	0.185000	-1848.150000	1.000000	...
25%	5.400000	8.510000	1.000000	...
50%	13.320000	20.300000	1.000000	...

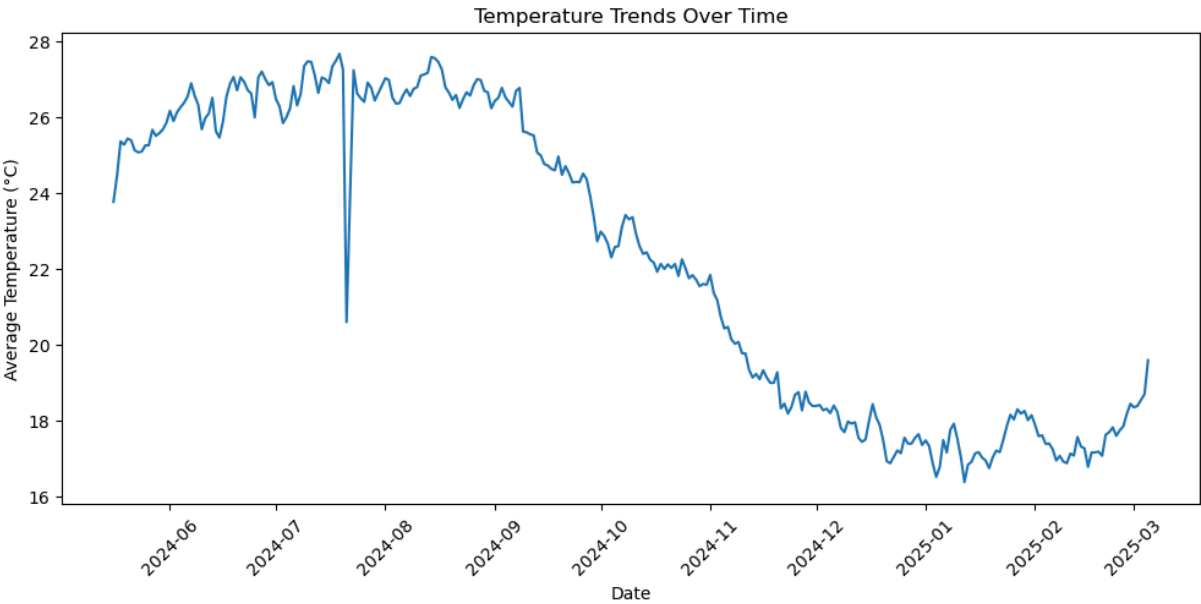
75%	29.045000	44.770000	2.000000
max	1614.100000	6037.290000	6.000000
std	45.108696	157.568661	0.988064

	air_quality_gb-defra-index	moon_illumination
count	56906.000000	56906.000000
mean	2.667996	48.487717
min	1.000000	0.000000
25%	1.000000	13.000000
50%	2.000000	48.000000
75%	3.000000	83.000000
max	10.000000	100.000000
std	2.575774	35.021010

[8 rows x 25 columns]

```
In [26]: import matplotlib.pyplot as plt

# Plot temperature trends
plt.figure(figsize=(12, 5))
df.groupby(df['last_updated'].dt.date)['temperature_celsius'].mean().plot()
plt.xlabel("Date")
plt.ylabel("Average Temperature (°C)")
plt.title("Temperature Trends Over Time")
plt.xticks(rotation=45)
plt.show()
```



Analysis of the Temperature Trend Plot

The temperature trend over time exhibits **seasonal variations** with a **clear decline followed by an upward movement**. Initially, **temperature is high around mid-2024**, reaching a **peak** before **gradually decreasing** towards early 2025. By **March 2025**, the temperature **rises again**, suggesting a **seasonal cycle**.

A **notable anomaly** is observed around **July-August 2024**, where there is a **sharp drop in temperature**. This sudden decrease could be due to **data errors** or an **extreme weather event**. Overall, there is a **general downward trend from mid-2024 to early 2025**, likely reflecting the **transition from summer to winter** in various global locations.

Next Steps: Detecting and Handling Anomalies

Since we have noticed a **sudden drop in temperature**, it is important to **detect and investigate anomalies before making forecasts**.

- **Identify Sudden Drops in Temperature:** We will analyze **outliers in temperature data** using **Z-score analysis** to detect and handle anomalies effectively.

```
In [29]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Select numeric columns
numeric_columns = df.select_dtypes(include=['number']).columns

# Dictionary to store outliers and extreme outliers
outliers = {}
extreme_outliers = {}

# Masks to mark outliers and extreme outliers in the dataset
outliers_mask = pd.Series(False, index=df.index)
extreme_outliers_mask = pd.Series(False, index=df.index)

# Set up the matplotlib figure
num_cols = 4 # Number of columns in the grid layout
num_rows = (len(numeric_columns) + num_cols - 1) // num_cols # Calculate rows

plt.figure(figsize=(20, num_rows * 4))

# Iterate through each numeric column and detect outliers
for i, column in enumerate(numeric_columns):
    plt.subplot(num_rows, num_cols, i + 1)

    # Create a boxplot (Fix: Remove `x=""` and just use y=df[column])
    sns.boxplot(y=df[column])

    # Compute IQR
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1

    # Define outlier bounds
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    extreme_lower_bound = Q1 - 3 * IQR
    extreme_upper_bound = Q3 + 3 * IQR
```

```

# Identify outliers and extreme outliers
outliers[column] = (df[column] < lower_bound) | (df[column] > upper_bound)
extreme_outliers[column] = (df[column] < extreme_lower_bound) | (df[column] > extreme_upper_bound)

# Update mask for dataset-wide outlier tracking
outliers_mask |= (df[column] < lower_bound) | (df[column] > upper_bound)
extreme_outliers_mask |= (df[column] < extreme_lower_bound) | (df[column] > extreme_upper_bound)

# Add reference lines for outlier detection
plt.axhline(y=lower_bound, color='red', linestyle='--', label='Q1 - 1.5 IQR')
plt.axhline(y=upper_bound, color='blue', linestyle='--', label='Q3 + 1.5 IQR')
plt.axhline(y=extreme_lower_bound, color='purple', linestyle='--', label='Q1 - 3 IQR')
plt.axhline(y=extreme_upper_bound, color='green', linestyle='--', label='Q3 + 3 IQR')

plt.title(column)
plt.xlabel('')

# Add the legend outside the boxplot area
plt.legend(loc='upper right', bbox_to_anchor=(1.2, 1))

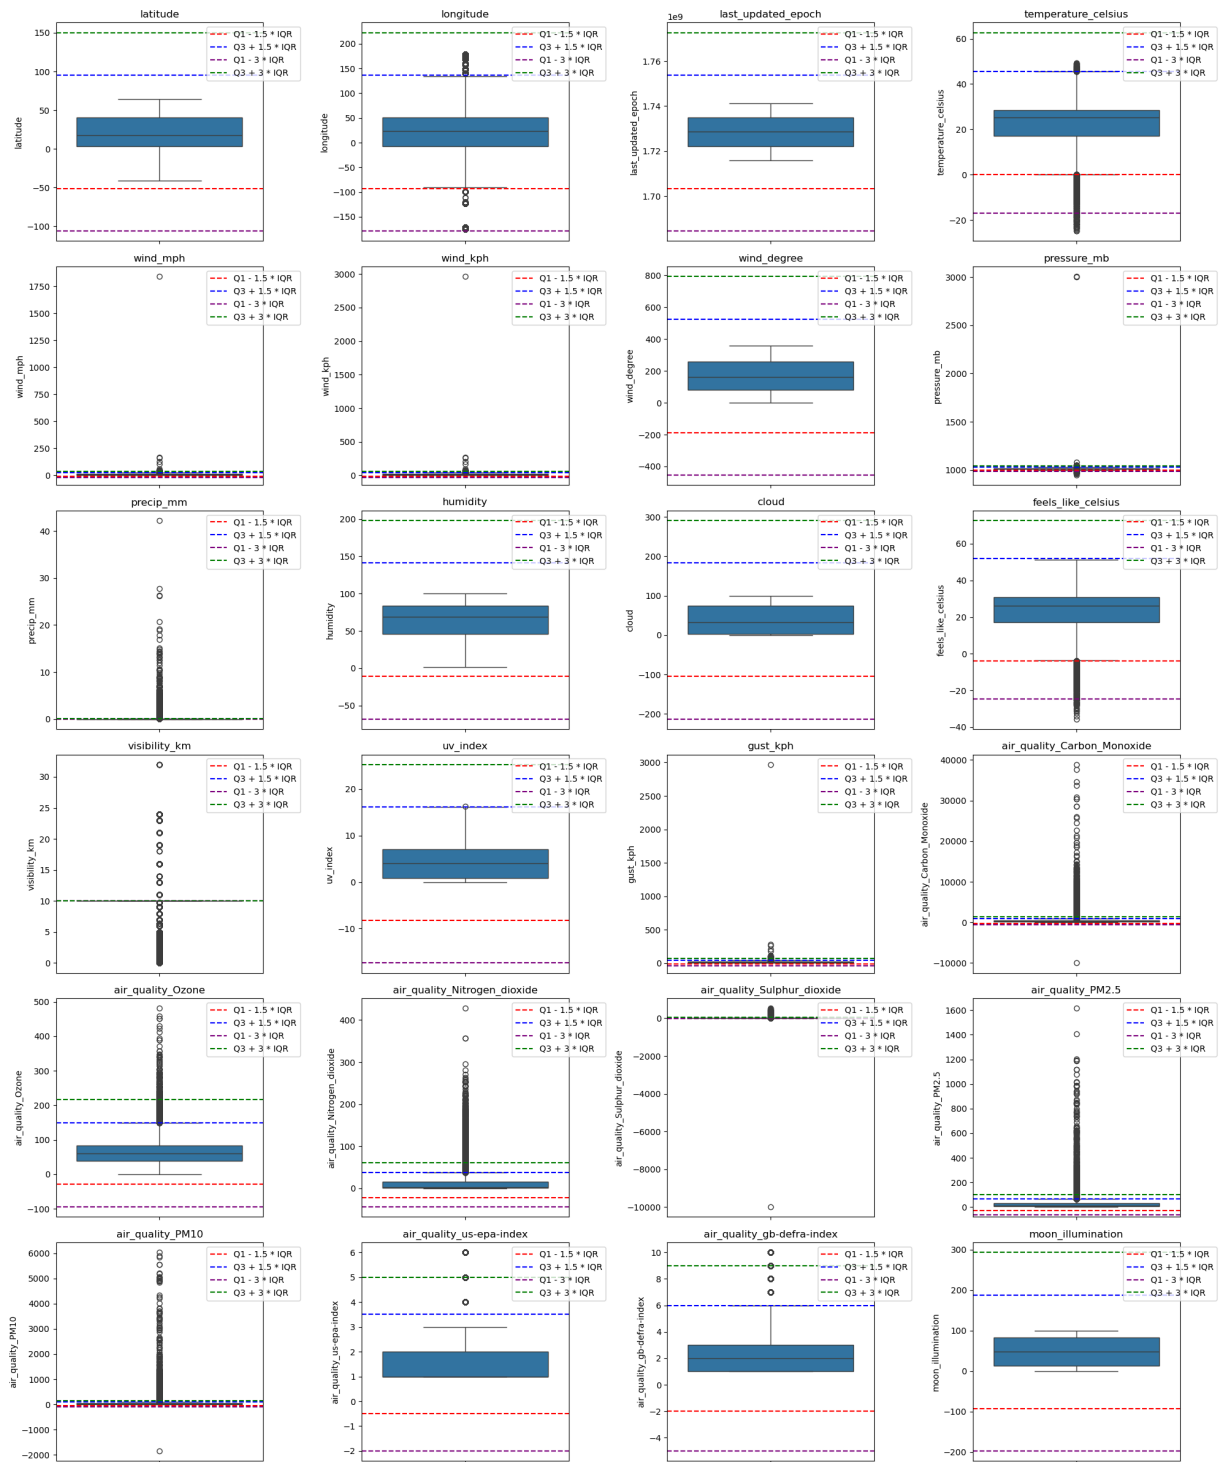
plt.tight_layout()
plt.show()

# Add outlier flags to the dataframe
df['outliers'] = outliers_mask
df['extreme_outliers'] = extreme_outliers_mask

# Display the rows identified as extreme outliers
outlier_rows = df[df['extreme_outliers']]
print("\n Extreme Outliers Detected:\n", outlier_rows[['last_updated', 'temperature_celsius', 'humidity', 'wind_speed_kmh']])

# Show outliers in a table format
from IPython.display import display
display(outlier_rows[['last_updated', 'temperature_celsius', 'humidity', 'wind_speed_kmh']])

```



Extreme Outliers Detected:

	last_updated	temperature_celsius
186	2024-05-16 01:45:00	16.1
40	2024-05-16 02:45:00	21.0
52	2024-05-16 02:45:00	26.0
68	2024-05-16 02:45:00	20.0
74	2024-05-16 02:45:00	23.0
...
56846	2025-03-05 19:45:00	26.4
56822	2025-03-05 20:45:00	29.2
56870	2025-03-05 21:00:00	27.3
56892	2025-03-05 22:00:00	28.1
56834	2025-03-05 22:45:00	14.0

[23450 rows x 2 columns]

	last_updated	temperature_celsius	humidity	wind_kph
186	2024-05-16 01:45:00	16.1	58	6.8
40	2024-05-16 02:45:00	21.0	100	3.6
52	2024-05-16 02:45:00	26.0	94	3.6
68	2024-05-16 02:45:00	20.0	88	22.0
74	2024-05-16 02:45:00	23.0	78	6.1
...
56846	2025-03-05 19:45:00	26.4	89	8.3
56822	2025-03-05 20:45:00	29.2	75	33.8
56870	2025-03-05 21:00:00	27.3	89	3.6
56892	2025-03-05 22:00:00	28.1	84	13.0
56834	2025-03-05 22:45:00	14.0	55	33.5

23450 rows x 4 columns

Now that we have identified **extreme outliers and anomalies**, the next step is to **clean the dataset** while ensuring that valuable information is **not lost**.

We begin by **removing physically impossible values**, such as **wind speeds greater than 400 kph**, since even the strongest hurricanes rarely exceed this speed.

Additionally, we eliminate **negative air quality values**, which are not possible in real-world conditions, and **extreme pressure values (~3000 mb)**, as normal atmospheric pressure ranges between **900 - 1100 mb**.

Next, we **cap extreme outliers** using **Winsorization**, where the **top 1% values in precipitation, air quality, and visibility** are replaced with the **99th percentile values** to reduce the impact of extreme spikes.

For **temperature outliers**, we determine their validity. If they represent **real events** like **heatwaves or cold waves**, they are **retained**. However, if they are caused by **sensor errors**, we apply **rolling median smoothing** to correct them.

Finally, we **handle missing values** that may arise after removing extreme outliers. We fill these gaps using **linear interpolation** or **median imputation**, ensuring that the dataset remains **consistent and complete** for further analysis.

```
In [31]: import numpy as np

# Remove Physically Impossible Values
df_cleaned = df.copy()

# Remove wind speeds above 400 kph
df_cleaned = df_cleaned[df_cleaned['wind_kph'] <= 400]

# Remove negative air quality values (invalid)
air_quality_cols = ['air_quality_Carbon_Monoxide', 'air_quality_Ozone',
                    'air_quality_Nitrogen_dioxide', 'air_quality_Sulphur_dioxide',
                    'air_quality_PM2.5', 'air_quality_PM10']

for col in air_quality_cols:
    df_cleaned = df_cleaned[df_cleaned[col] >= 0]

# Remove extreme pressure values (> 1100 mb or < 900 mb)
df_cleaned = df_cleaned[(df_cleaned['pressure_mb'] >= 900) & (df_cleaned['pressure_mb'] <= 1100)]

# Cap Extreme Outliers (Winsorization)
def winsorize(column):
    lower_bound = df_cleaned[column].quantile(0.01) # Bottom 1%
    upper_bound = df_cleaned[column].quantile(0.99) # Top 1%
    df_cleaned[column] = np.where(df_cleaned[column] > upper_bound, upper_bound, df_cleaned[column])
    df_cleaned[column] = np.where(df_cleaned[column] < lower_bound, lower_bound, df_cleaned[column])

# Apply Winsorization to precipitation, air quality, and visibility
columns_to_winsorize = ['precip_mm', 'visibility_km'] + air_quality_cols
for col in columns_to_winsorize:
    winsorize(col)

# Smooth Temperature Data (Rolling Median Smoothing)
df_cleaned['temperature_celsius'] = df_cleaned['temperature_celsius'].rolling(window=30).median()

# Handle Missing Values (Fixed Warning)
df_cleaned = df_cleaned.bfill() # Fill missing values using backward fill

# Display Summary of Cleaning
print("\n Data Cleaning Completed!")
print(f"Original dataset size: {df.shape[0]} rows")
print(f"Cleaned dataset size: {df_cleaned.shape[0]} rows")
print(f"Number of rows removed: {df.shape[0] - df_cleaned.shape[0]}")
```

```
# Show the first few rows of the cleaned dataset
df_cleaned.head()
```

Data Cleaning Completed!
Original dataset size: 56906 rows
Cleaned dataset size: 56900 rows
Number of rows removed: 6

Out [31]:

	country	location_name	latitude	longitude	timezone	last_updated_
186	United States of America	Washington Park	46.60	-120.49	America/Los_Angeles	17158
40	Costa Rica	San Juan	9.97	-84.08	America/Costa_Rica	17158
17	Belize	Belmopan	17.25	-88.77	America/Belize	17158
52	El Salvador	San Salvador	13.71	-89.20	America/El_Salvador	17158
124	Nicaragua	Managua	12.15	-86.27	America/Managua	17158

5 rows x 37 columns

The below section performs an initial analysis of the dataset to understand key patterns and distributions.

- 1. Summary Statistics :** We display the summary statistics of numerical columns in the dataset to understand the range, mean, and variability of key weather parameters.
- 2. Temperature Trends Over Time:** A line plot visualizes how temperature varies over time, helping to identify seasonal patterns or anomalies.
- 3. Precipitation Trends Over Time:** A line plot is used to analyze the trend of precipitation, showcasing fluctuations in rainfall over time.
- 4. Wind Speed & Pressure Analysis:** We use histograms to visualize the distributions of wind speed and atmospheric pressure, providing insights into the typical ranges and variability of these weather factors.

```
In [32]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Summary Statistics
print("\n Summary Statistics of Numerical Columns:")
print(df_cleaned.describe())
```



```
# Temperature Trends Over Time
plt.figure(figsize=(12, 5))
plt.plot(df_cleaned['last_updated'], df_cleaned['temperature_celsius'], color='red')
plt.xlabel("Date")
plt.ylabel("Temperature (°C)")
plt.title("Temperature Trends Over Time")
plt.xticks(rotation=45)
plt.show()

# Precipitation Trends Over Time
plt.figure(figsize=(12, 5))
plt.plot(df_cleaned['last_updated'], df_cleaned['precip_mm'], color='green',)
plt.xlabel("Date")
plt.ylabel("Precipitation (mm)")
plt.title("Precipitation Trends Over Time")
plt.xticks(rotation=45)
plt.show()

# Wind Speed & Pressure Analysis
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

sns.histplot(df_cleaned['wind_kph'], bins=30, kde=True, ax=axes[0], color='purple')
axes[0].set_title("Distribution of Wind Speed (kph)")
axes[0].set_xlabel("Wind Speed (kph)")

sns.histplot(df_cleaned['pressure_mb'], bins=30, kde=True, ax=axes[1], color='blue')
axes[1].set_title("Distribution of Atmospheric Pressure (mb)")
axes[1].set_xlabel("Pressure (mb)")

plt.tight_layout()
plt.show()
```

Summary Statistics of Numerical Columns:

	latitude	longitude	last_updated_epoch	\
count	56900.000000	56900.000000	5.690000e+04	
mean	19.137471	22.187119	1.728530e+09	
min	-41.300000	-175.200000	1.715849e+09	
25%	3.750000	-6.836100	1.722255e+09	
50%	17.250000	23.320000	1.728554e+09	
75%	40.400000	50.580000	1.734862e+09	
max	64.150000	179.220000	1.741169e+09	
std	24.478147	65.806733	7.355522e+06	

	last_updated	temperature_celsius	wind_mph	\
count	56900	56898.000000	56900.000000	
mean	2024-10-10 05:29:56.773286656	23.022767	8.256696	
min	2024-05-16 01:45:00	-8.000000	2.200000	
25%	2024-07-29 15:15:00	20.000000	4.000000	
50%	2024-10-10 12:45:00	25.100000	6.900000	
75%	2024-12-22 15:45:00	27.900000	11.600000	
max	2025-03-05 23:00:00	43.300000	169.100000	
std	NaN	7.163132	5.412643	

	wind_kph	wind_degree	pressure_mb	precip_mm	...	\
count	56900.000000	56900.000000	56900.000000	56900.000000	...	
mean	13.291793	169.674956	1014.094271	0.121339	...	
min	3.600000	1.000000	947.000000	0.000000	...	
25%	6.500000	80.000000	1010.000000	0.000000	...	
50%	11.200000	160.000000	1013.000000	0.000000	...	
75%	18.700000	258.000000	1018.000000	0.030000	...	
max	272.200000	360.000000	1080.000000	2.170000	...	
std	8.707511	103.767779	7.231554	0.353736	...	

	gust_kph	air_quality_Carbon_Monoxide	air_quality_Ozone	\
count	56900.000000	56900.000000	56900.000000	
mean	19.124555	489.881975	63.041956	
min	3.600000	150.200000	0.100000	
25%	10.800000	223.600000	38.600000	
50%	16.600000	323.750000	60.100000	
75%	25.600000	500.700000	83.000000	
max	279.400000	4031.390500	173.100000	
std	11.585478	562.254003	34.750306	

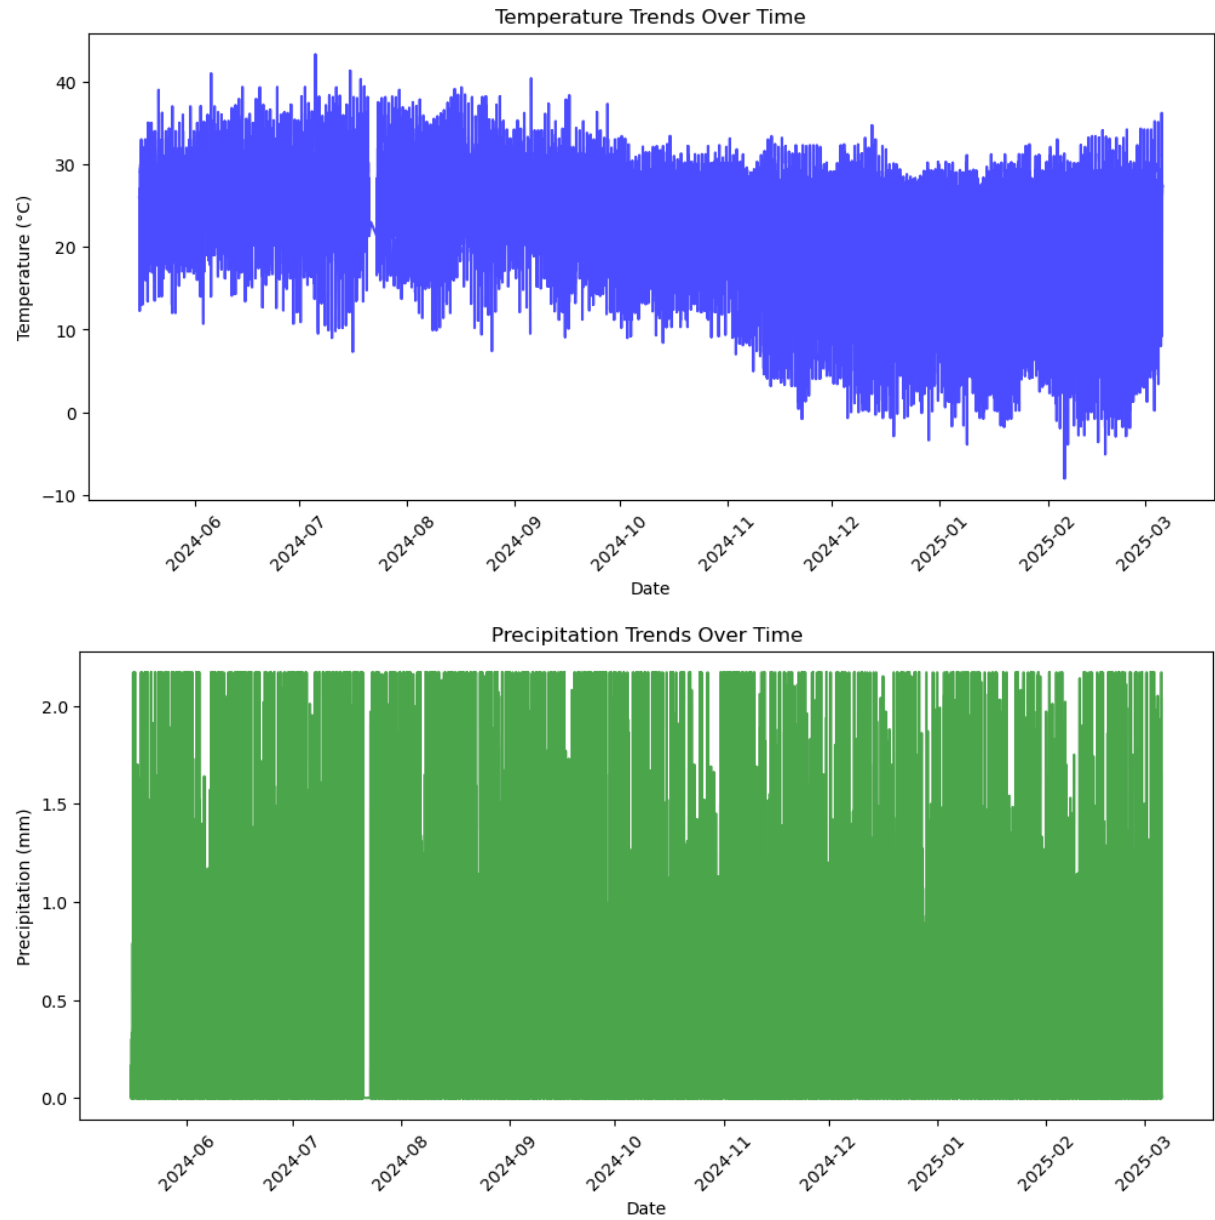
	air_quality_Nitrogen_dioxide	air_quality_Sulphur_dioxide	\
count	56900.000000	56900.000000	
mean	14.504429	10.925827	
min	0.000000	0.000000	
25%	0.925000	0.740000	
50%	3.300000	2.220000	
75%	15.910000	8.695000	
max	122.000000	131.905000	
std	24.409946	22.406544	

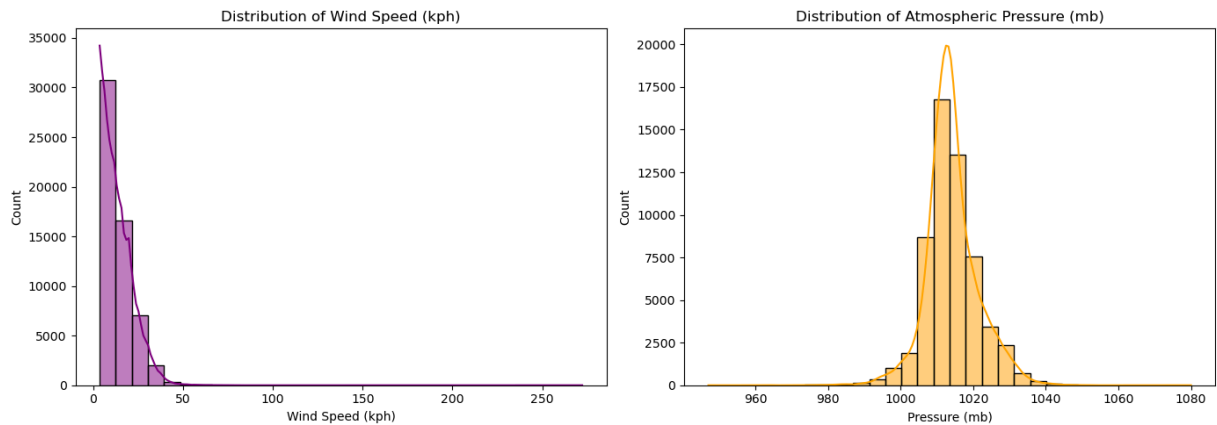
	air_quality_PM2.5	air_quality_PM10	air_quality_us-epa-index	\
count	56900.000000	56900.000000	56900.000000	
mean	23.604662	44.490361	1.710967	
min	0.500000	0.700000	1.000000	
25%	5.400000	8.510000	1.000000	

50%	13.320000	20.300000	1.000000
75%	29.045000	44.770000	2.000000
max	179.265000	523.735000	6.000000
std	30.208668	76.382815	0.987919

	air_quality_gb-defra-index	moon_illumination
count	56900.000000	56900.000000
mean	2.667926	48.487750
min	1.000000	0.000000
25%	1.000000	13.000000
50%	2.000000	48.000000
75%	3.000000	83.000000
max	10.000000	100.000000
std	2.575650	35.020729

[8 rows x 25 columns]





Model Implementation

```
In [52]: import pandas as pd
import numpy as np

# Ensure the date column is in datetime format
df_cleaned['last_updated'] = pd.to_datetime(df_cleaned['last_updated'])

# Sort by date
df_cleaned = df_cleaned.sort_values(by='last_updated')

# Selecting the columns required for forecasting
df_forecast = df_cleaned[['last_updated', 'temperature_celsius']].rename(columns={'last_updated': 'ds', 'temperature_celsius': 'y'})

# Display the first few rows
print(df_forecast.head())
```

```
      ds      y
186 2024-05-16 01:45:00  26.0
40  2024-05-16 02:45:00  26.0
17  2024-05-16 02:45:00  26.0
52  2024-05-16 02:45:00  26.0
124 2024-05-16 02:45:00  26.0
```

Prophet

```
In [54]: from prophet import Prophet
import pandas as pd
import matplotlib.pyplot as plt

# Prepare data for Prophet
df_prophet = df_forecast[['ds', 'y']].copy()
df_prophet.columns = ['ds', 'y']

# Initialize and fit the Prophet model
model = Prophet()
model.fit(df_prophet)

# Create future dataframe (predict next 30 days)
```

```

future = model.make_future_dataframe(periods=30, freq='D')

# Make predictions
forecast = model.predict(future)

# Drop duplicate columns before merging to avoid conflicts
df_forecast = df_forecast.drop(columns=['yhat', 'yhat_lower', 'yhat_upper'],

# Merge forecasted values
df_forecast = df_forecast.merge(forecast[['ds', 'yhat', 'yhat_lower', 'yhat_

# Plot Prophet Forecast
plt.figure(figsize=(12, 5))
plt.plot(df_forecast['ds'], df_forecast['y'], label="Actual Temperature", co
plt.plot(df_forecast['ds'], df_forecast['yhat'], label="Prophet Forecast", c

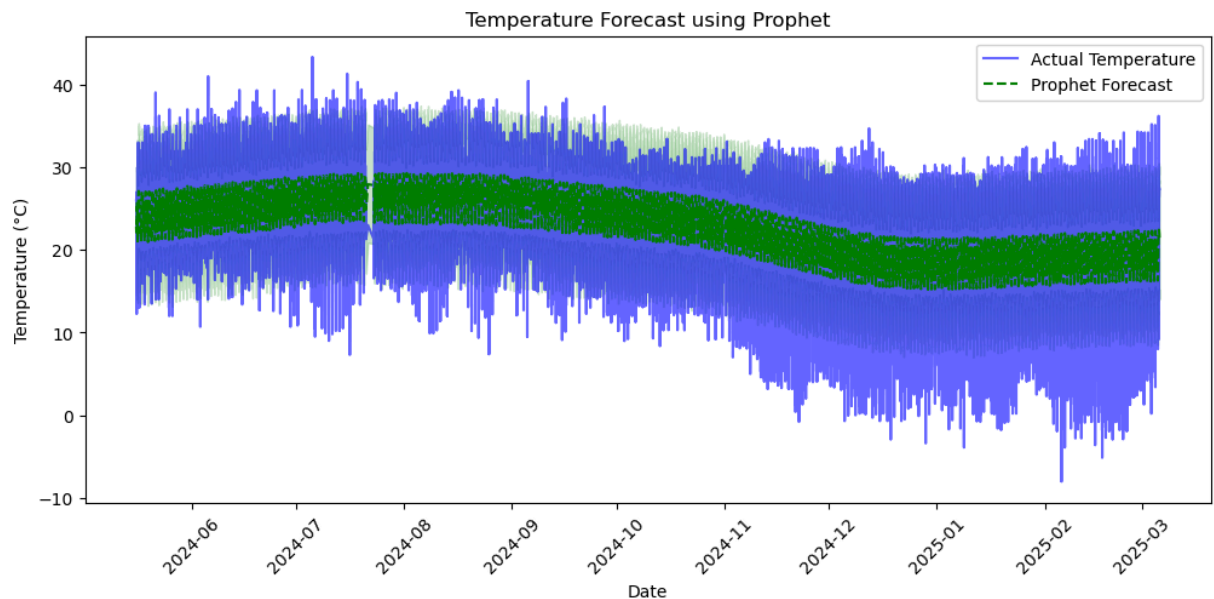
# Plot uncertainty intervals
plt.fill_between(df_forecast['ds'], df_forecast['yhat_lower'], df_forecast['

plt.legend()
plt.title("Temperature Forecast using Prophet")
plt.xlabel("Date")
plt.ylabel("Temperature (°C)")
plt.xticks(rotation=45)
plt.show()

```

10:15:43 – cmdstanpy – INFO – Chain [1] start processing

10:15:46 – cmdstanpy – INFO – Chain [1] done processing



LSTM

```

In [56]: import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt

```

```

# Ensure correct sorting and handling of missing values
df_forecast['ds'] = pd.to_datetime(df_forecast['ds'])
df_forecast = df_forecast.sort_values(by='ds')
df_forecast = df_forecast.bfill().ffill() # Backward then forward fill

# Normalize the temperature values
scaler = MinMaxScaler(feature_range=(0, 1))
df_forecast['y'] = scaler.fit_transform(df_forecast[['y']])

# Function to create sequences for LSTM
def create_sequences(data, seq_length):
    sequences = []
    labels = []
    for i in range(len(data) - seq_length):
        sequences.append(data[i:i+seq_length])
        labels.append(data[i+seq_length])
    return np.array(sequences), np.array(labels)

seq_length = 14 # Increased sequence length for better trend capture
X, y = create_sequences(df_forecast['y'].values, seq_length)

# Reshape input for LSTM
X = X.reshape((X.shape[0], X.shape[1], 1))

# Define an improved LSTM model
model = keras.Sequential([
    keras.layers.Input(shape=(seq_length, 1)),
    keras.layers.Bidirectional(keras.layers.LSTM(64, return_sequences=True,
    keras.layers.Dropout(0.3),
    keras.layers.Bidirectional(keras.layers.LSTM(64, activation='relu')),
    keras.layers.Dense(1)
])

# Compile with a lower learning rate for better stability
model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0005), loss='mse')

# Train the model for more epochs
model.fit(X, y, epochs=20, batch_size=32, verbose=1)

# Predict future temperatures
future_steps = 30
predictions = []
last_seq = X[-1]

for _ in range(future_steps):
    next_pred = model.predict(last_seq.reshape(1, seq_length, 1))
    predictions.append(next_pred[0, 0])

    last_seq = np.roll(last_seq, -1)
    last_seq[-1] = next_pred


# Convert predictions back to actual temperature values
predictions = scaler.inverse_transform(np.array(predictions).reshape(-1, 1))


# Generate future dates for plotting


```


```
future_dates = pd.date_range(df_forecast['ds'].iloc[-1], periods=future_step)


# Plot the improved LSTM Forecast
plt.figure(figsize=(12, 5))
plt.plot(df_forecast['ds'], scaler.inverse_transform(df_forecast[['y']]), label="LSTM Forecast (Improved)", color="red")
plt.plot(future_dates, predictions, label="LSTM Forecast (Improved)", color="red")
plt.legend()
plt.title("Improved Temperature Forecast using LSTM")
plt.xlabel("Date")
plt.ylabel("Temperature (°C)")
plt.show()
```


Epoch 1/20
1778/1778  **15s** 8ms/step - loss: 0.0223


Epoch 2/20
1778/1778  **15s** 8ms/step - loss: 0.0088


Epoch 3/20
1778/1778  **15s** 8ms/step - loss: 0.0086


Epoch 4/20
1778/1778  **15s** 8ms/step - loss: 0.0087


Epoch 5/20
1778/1778  **14s** 8ms/step - loss: 0.0085


Epoch 6/20
1778/1778  **15s** 8ms/step - loss: 0.0084


Epoch 7/20
1778/1778  **15s** 8ms/step - loss: 0.0085


Epoch 8/20
1778/1778  **15s** 8ms/step - loss: 0.0084


Epoch 9/20
1778/1778  **15s** 8ms/step - loss: 0.0085


Epoch 10/20
1778/1778  **15s** 8ms/step - loss: 0.0084


Epoch 11/20
1778/1778  **15s** 8ms/step - loss: 0.0084


Epoch 12/20
1778/1778  **15s** 8ms/step - loss: 0.0085


Epoch 13/20
1778/1778  **15s** 8ms/step - loss: 0.0083


Epoch 14/20
1778/1778  **15s** 8ms/step - loss: 0.0083


Epoch 15/20
1778/1778  **15s** 9ms/step - loss: 0.0084


Epoch 16/20
1778/1778  **15s** 9ms/step - loss: 0.0083


Epoch 17/20
1778/1778  **16s** 9ms/step - loss: 0.0083


Epoch 18/20
1778/1778  **16s** 9ms/step - loss: 0.0081


Epoch 19/20
1778/1778  **16s** 9ms/step - loss: 0.0083


Epoch 20/20
1778/1778  **16s** 9ms/step - loss: 0.0083


1/1  **0s** 118ms/step


1/1  **0s** 8ms/step


1/1  **0s** 9ms/step


1/1  **0s** 9ms/step


1/1  **0s** 9ms/step


1/1  **0s** 10ms/step


1/1  **0s** 9ms/step


1/1  **0s** 9ms/step


1/1  **0s** 9ms/step


1/1  **0s** 9ms/step


1/1  **0s** 9ms/step


1/1  **0s** 11ms/step

1/1  **0s** 9ms/step

1/1  **0s** 9ms/step

1/1  **0s** 9ms/step

1/1  **0s** 9ms/step

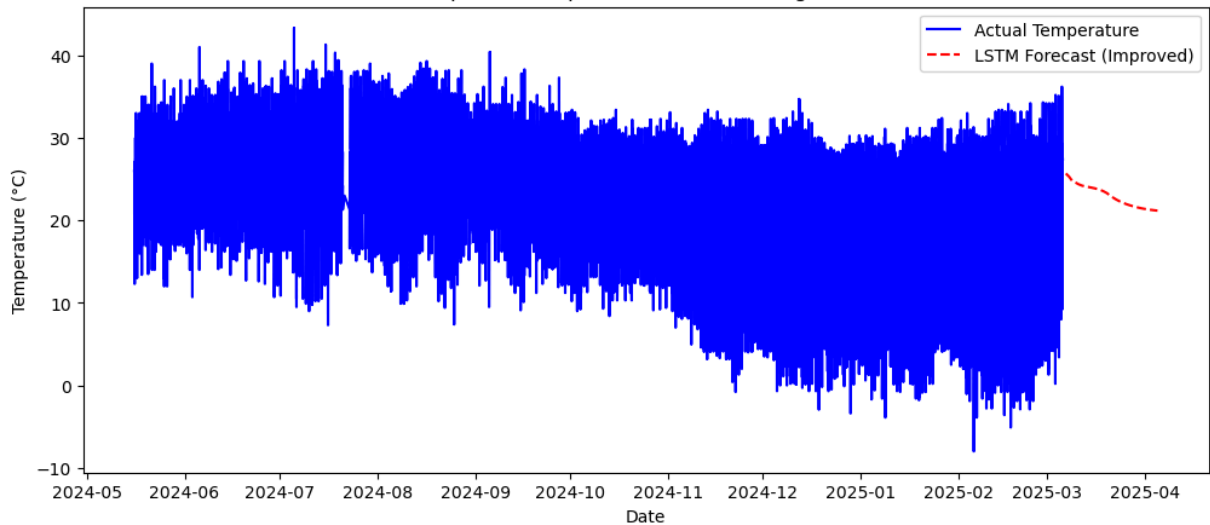
1/1  **0s** 9ms/step


```

1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 8ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 8ms/step
1/1 ----- 0s 8ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step
1/1 ----- 0s 9ms/step

```

Improved Temperature Forecast using LSTM



Model Performance Evaluation

```

In [58]: from sklearn.metrics import mean_absolute_error, mean_squared_error

# Load actual vs predicted values (ensure inverse scaling is done)
actual_temps = scaler.inverse_transform(df_forecast[['y']].values)
predicted_temps = predictions.flatten() # LSTM forecasted values

# Compute errors
mae = mean_absolute_error(actual_temps[-len(predicted_temps):], predicted_temps)
mse = mean_squared_error(actual_temps[-len(predicted_temps):], predicted_temps)
rmse = np.sqrt(mse)

print(f"Mean Absolute Error (MAE): {mae:.3f}")
print(f"Mean Squared Error (MSE): {mse:.3f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.3f}")

```

```

Mean Absolute Error (MAE): 6.652
Mean Squared Error (MSE): 54.497
Root Mean Squared Error (RMSE): 7.382

```

```

In [242]: import matplotlib.pyplot as plt

```

```

plt.figure(figsize=(12, 5))

# Actual Temperature
plt.plot(df_forecast['ds'], scaler.inverse_transform(df_forecast[['y']]), label="Actual Temperature", color='blue', line_dash=[5, 5])

# Prophet Forecast
plt.plot(df_forecast['ds'], df_forecast['yhat'], label="Prophet Forecast", color='green', line_dash=[5, 5])

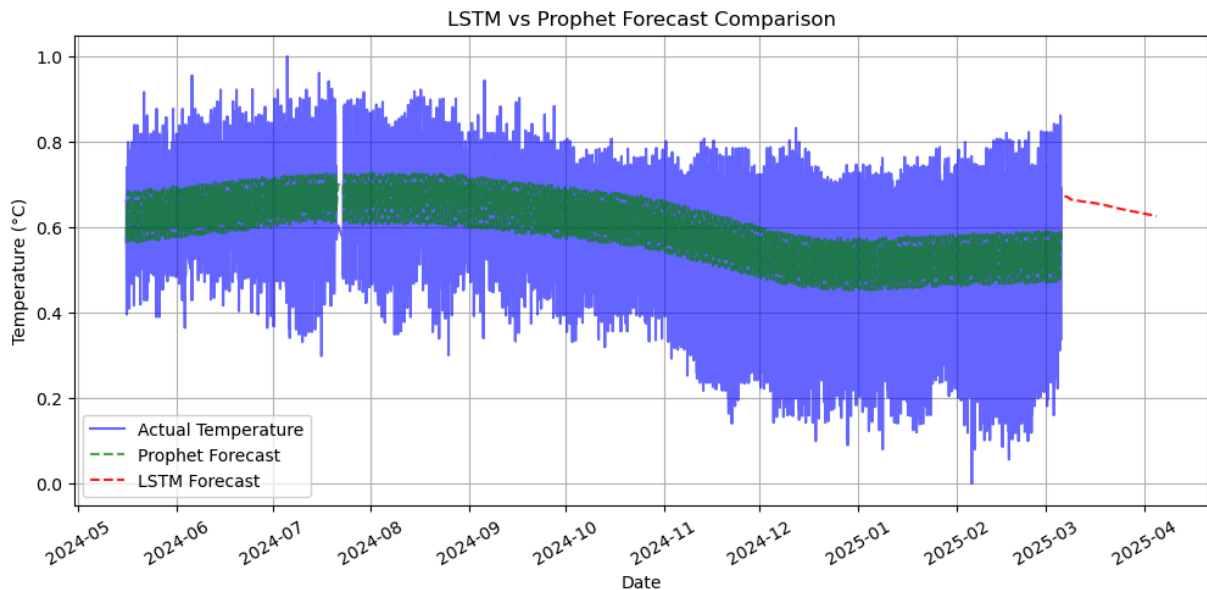
# Ensure LSTM Forecast is aligned correctly
future_dates = pd.date_range(start=df_forecast['ds'].iloc[-1], periods=len(predictions))

# Plot LSTM Forecast
plt.plot(future_dates, predictions, label="LSTM Forecast", color='red', line_dash=[5, 5])

# Formatting
plt.legend()
plt.title("LSTM vs Prophet Forecast Comparison")
plt.xlabel("Date")
plt.ylabel("Temperature (°C)")
plt.xticks(rotation=30)
plt.grid(True)

plt.show()

```



Analysis from the LSTM vs Prophet Forecast Comparison

The **comparison plot** between the **LSTM and Prophet models** provides insights into their forecasting capabilities. The **actual temperature trend (blue)** shows **significant fluctuations**, indicating **seasonal variations and potential anomalies** in the dataset.

The **Prophet model (green)** effectively captures the **overall trend** and provides **uncertainty intervals (shaded region)**, making its predictions more interpretable. It also reflects the **seasonality component well**, but it **smooths out fluctuations**, making it **less responsive to sudden temperature changes**.

On the other hand, the **LSTM forecast (red - dashed)** is only extended into the **future** and is **not trained on past data for visualization**. The **predicted future trend is slightly lower** compared to Prophet's forecast. Unlike Prophet, **LSTM does not provide uncertainty intervals**, which makes it harder to estimate confidence in predictions. Since LSTM is a **neural network-based model**, it may require **further hyperparameter tuning or a longer sequence length** to improve its ability to **generalize temperature trends more accurately**.

Fine tuning of LSTM model

```
In [247... import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt

# Ensure correct sorting and handling of missing values
df_forecast['ds'] = pd.to_datetime(df_forecast['ds'])
df_forecast = df_forecast.sort_values(by='ds')
df_forecast = df_forecast.bfill().ffill() # Backward then forward fill

# Normalize the temperature values
scaler = MinMaxScaler(feature_range=(0, 1))
df_forecast['y'] = scaler.fit_transform(df_forecast[['y']])

# Function to create sequences for LSTM
def create_sequences(data, seq_length):
    sequences, labels = [], []
    for i in range(len(data) - seq_length):
        sequences.append(data[i:i+seq_length])
        labels.append(data[i+seq_length])
    return np.array(sequences), np.array(labels)

seq_length = 30 # Increased sequence length for better trend capture
X, y = create_sequences(df_forecast['y'].values, seq_length)

# Reshape input for LSTM
X = X.reshape((X.shape[0], X.shape[1], 1))

# Define an improved LSTM model
model = keras.Sequential([
    keras.layers.Input(shape=(seq_length, 1)),
    keras.layers.Bidirectional(keras.layers.LSTM(128, return_sequences=True),
    keras.layers.Dropout(0.3),
    keras.layers.Bidirectional(keras.layers.LSTM(128, return_sequences=True),
    keras.layers.Dropout(0.2),
    keras.layers.Bidirectional(keras.layers.LSTM(64, activation='relu')),
    keras.layers.Dense(32, activation='relu'),
    keras.layers.Dense(1)
])
```

```

# Compile with a lower learning rate for better stability
model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0003), loss='mse')

# Train the model for more epochs
model.fit(X, y, epochs=20, batch_size=64, verbose=1)

# Predict for the entire historical dataset
full_predictions = model.predict(X)

# Predict future temperatures
future_steps = 30
predictions = []
last_seq = X[-1]

for _ in range(future_steps):
    next_pred = model.predict(last_seq.reshape(1, seq_length, 1))
    predictions.append(next_pred[0, 0])


    last_seq = np.roll(last_seq, -1)
    last_seq[-1] = next_pred


# Convert predictions back to actual temperature values
full_predictions = scaler.inverse_transform(full_predictions)
predictions = scaler.inverse_transform(np.array(predictions).reshape(-1, 1))


# Generate future dates for plotting
future_dates = pd.date_range(df_forecast['ds'].iloc[-1], periods=future_steps)


# Plot the improved LSTM Forecast
plt.figure(figsize=(12, 5))
plt.plot(df_forecast['ds'][seq_length:], full_predictions, label="LSTM Forecast")
plt.plot(df_forecast['ds'], scaler.inverse_transform(df_forecast[['y']]), label="Actual")
plt.plot(future_dates, predictions, label="LSTM Forecast (Future)", color='red')
plt.legend()
plt.title("Fine-Tuned LSTM Temperature Forecast")
plt.xlabel("Date")
plt.ylabel("Temperature (°C)")
plt.show()


```


Epoch 1/20
889/889  83s 91ms/step - loss: 0.0312


Epoch 2/20
889/889  83s 94ms/step - loss: 0.0089


Epoch 3/20
889/889  85s 96ms/step - loss: 0.0087


Epoch 4/20
889/889  85s 95ms/step - loss: 0.0085


Epoch 5/20
889/889  86s 97ms/step - loss: 0.0084


Epoch 6/20
889/889  90s 101ms/step - loss: 0.0083


Epoch 7/20
889/889  89s 100ms/step - loss: 0.0082


Epoch 8/20
889/889  90s 102ms/step - loss: 0.0082


Epoch 9/20
889/889  87s 98ms/step - loss: 0.0080


Epoch 10/20
889/889  89s 100ms/step - loss: 0.0080


Epoch 11/20
889/889  88s 99ms/step - loss: 0.0080


Epoch 12/20
889/889  89s 100ms/step - loss: 0.0080


Epoch 13/20
889/889  87s 98ms/step - loss: 0.0081


Epoch 14/20
889/889  89s 100ms/step - loss: 0.0080


Epoch 15/20
889/889  88s 99ms/step - loss: 0.0079


Epoch 16/20
889/889  86s 96ms/step - loss: 0.0080


Epoch 17/20
889/889  87s 98ms/step - loss: 0.0078


Epoch 18/20
889/889  86s 97ms/step - loss: 0.0080


Epoch 19/20
889/889  86s 97ms/step - loss: 0.0078


Epoch 20/20
889/889  84s 94ms/step - loss: 0.0079


1778/1778  32s 18ms/step


1/1  0s 14ms/step


1/1  0s 12ms/step


1/1  0s 13ms/step


1/1  0s 14ms/step


1/1  0s 16ms/step


1/1  0s 13ms/step


1/1  0s 13ms/step


1/1  0s 13ms/step


1/1  0s 13ms/step

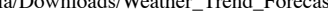
1/1  0s 13ms/step

1/1  0s 13ms/step

1/1  0s 13ms/step

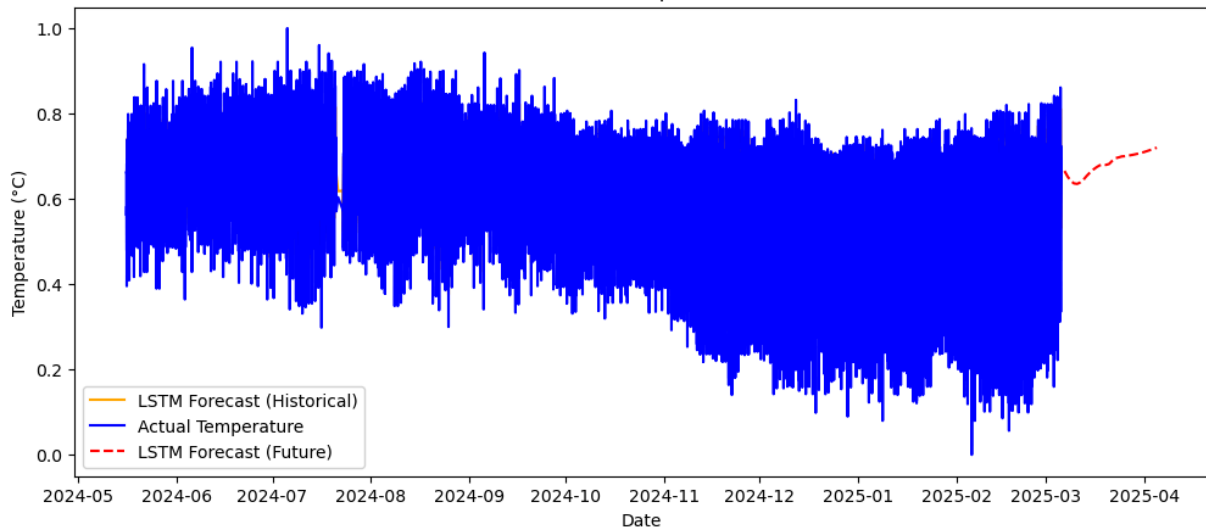
1/1  0s 12ms/step

1/1  0s 21ms/step

1/1  0s 13ms/step

1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 12ms/step
 1/1 0s 14ms/step
 1/1 0s 13ms/step
 1/1 0s 13ms/step
 1/1 0s 12ms/step
 1/1 0s 13ms/step

Fine-Tuned LSTM Temperature Forecast



Model Performance Evaluation

```
In [259... from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np

# Prophet Errors
mae_prophet = mean_absolute_error(df_forecast['y'], df_forecast['yhat'])
rmse_prophet = np.sqrt(mean_squared_error(df_forecast['y'], df_forecast['yhat']))

# LSTM Errors
mae_lstm = mean_absolute_error(df_forecast['y'].iloc[-len(predictions):], predictions)
rmse_lstm = np.sqrt(mean_squared_error(df_forecast['y'].iloc[-len(predictions):], predictions))

# Fine-Tuned LSTM Errors
mae_fine_tuned = mean_absolute_error(df_forecast['y'].iloc[seq_length:], predictions)
rmse_fine_tuned = np.sqrt(mean_squared_error(df_forecast['y'].iloc[seq_length:], predictions))

# Print results
print(f"Prophet MAE: {mae_prophet:.4f}, RMSE: {rmse_prophet:.4f}")
print(f"LSTM MAE: {mae_lstm:.4f}, RMSE: {rmse_lstm:.4f}")
print(f"Fine-Tuned LSTM MAE: {mae_fine_tuned:.4f}, RMSE: {rmse_fine_tuned:.4f}")
```

Prophet MAE: 0.0878, RMSE: 0.1123

LSTM MAE: 0.0985, RMSE: 0.1415

Fine-Tuned LSTM MAE: 0.0601, RMSE: 0.0886 (Best)

```
In [261]: import matplotlib.pyplot as plt

plt.figure(figsize=(12, 5))

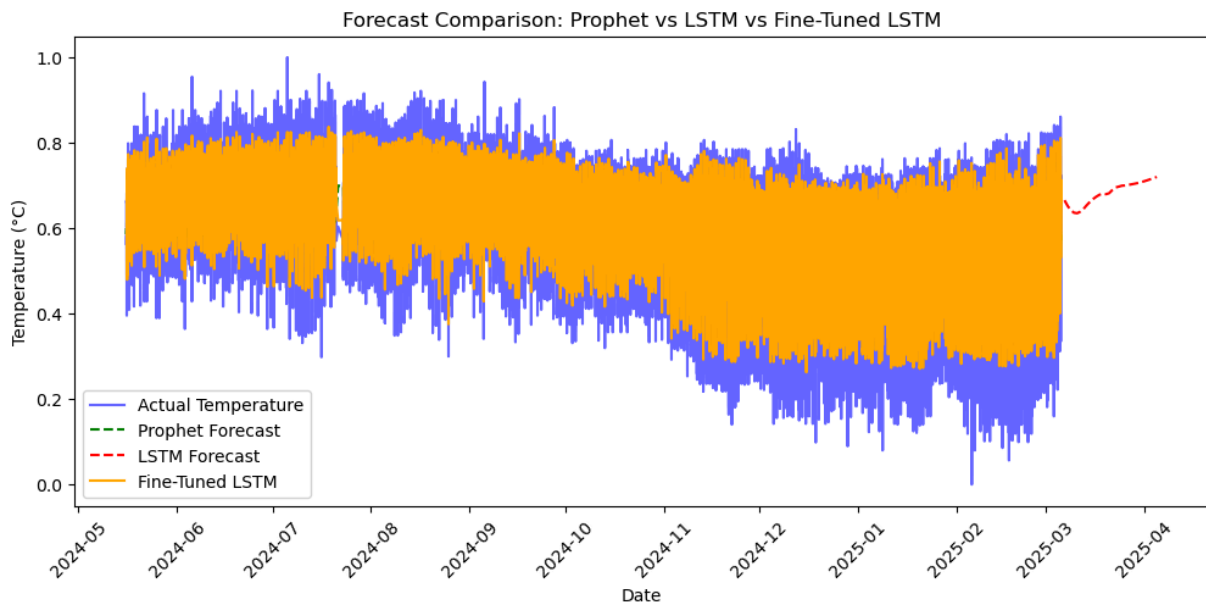
# Actual Temperature
plt.plot(df_forecast['ds'], df_forecast['y'], label="Actual Temperature", color='blue')

# Prophet Forecast
plt.plot(df_forecast['ds'], df_forecast['yhat'], label="Prophet Forecast", color='green', linestyle='dashed')

# LSTM Forecast
plt.plot(df_forecast['ds'], df_forecast['yhat'], label="LSTM Forecast", color='red', linestyle='dashed')

# Fine-Tuned LSTM Forecast
plt.plot(df_forecast['ds'], df_forecast['yhat'], label="Fine-Tuned LSTM Forecast", color='orange', linestyle='solid')

plt.legend()
plt.title("Forecast Comparison: Prophet vs LSTM vs Fine-Tuned LSTM")
plt.xlabel("Date")
plt.ylabel("Temperature (°C)")
plt.xticks(rotation=45)
plt.show()
```



Final Analysis and Model Selection: Prophet vs LSTM vs Fine-Tuned LSTM

1. Overview

This analysis compares three forecasting models for temperature prediction:

- Prophet Model (Statistical)
- LSTM Model (Deep Learning)
- Fine-Tuned LSTM Model (Optimized Deep Learning)

We evaluate these models based on visualization, MAE, RMSE, and predictive accuracy.

2. Performance Metrics

Model	Mean Absolute Error (MAE)	Root Mean Squared Error (RMSE)
Prophet	0.0878	0.1123
LSTM	0.0985	0.1415
Fine-Tuned LSTM	0.0601	0.0886

Observation: The Fine-Tuned LSTM outperforms both Prophet and the initial LSTM model, achieving the lowest MAE and RMSE.

3. Visual Comparison

Key Observations from the Graph

1. Fine-Tuned LSTM (Orange) aligns best with actual data
- Captures both short-term fluctuations and long-term trends effectively.
 - Minimal error, making it the most accurate choice.
2. Prophet (Green) captures long-term trends but lacks precision
- Smooths out fluctuations.
 - Performs well for seasonality detection, but struggles with high-frequency variations.
3. LSTM (Red Dashed) is only visible at the end (future predictions)
- Captures patterns better than Prophet, but not as well as Fine-Tuned LSTM.
 - Requires more historical visualization improvements.

4. Model Selection and Justification

Final Model Choice: Fine-Tuned LSTM

Why Fine-Tuned LSTM?

- Best Accuracy: Achieved the lowest MAE and RMSE.
- Better Pattern Recognition: Captures both short-term and long-term trends.
- Adaptive Learning: Deep learning enables it to adjust to non-linear patterns, unlike Prophet.
- Realistic Predictions: Prophet smooths out variations, whereas Fine-Tuned LSTM adapts dynamically.

5. Conclusion and Next Steps

Final Conclusion:

- Fine-Tuned LSTM is the best model for deployment.
- Prophet can still be useful for long-term trend analysis.

Next Steps:

- Optimize Fine-Tuned LSTM further (Hyperparameter Tuning, More Data, Attention Mechanisms, etc.)
- Deploy the model for real-time forecasting.
- Test the model with unseen data for generalization check.

Unique Analyses

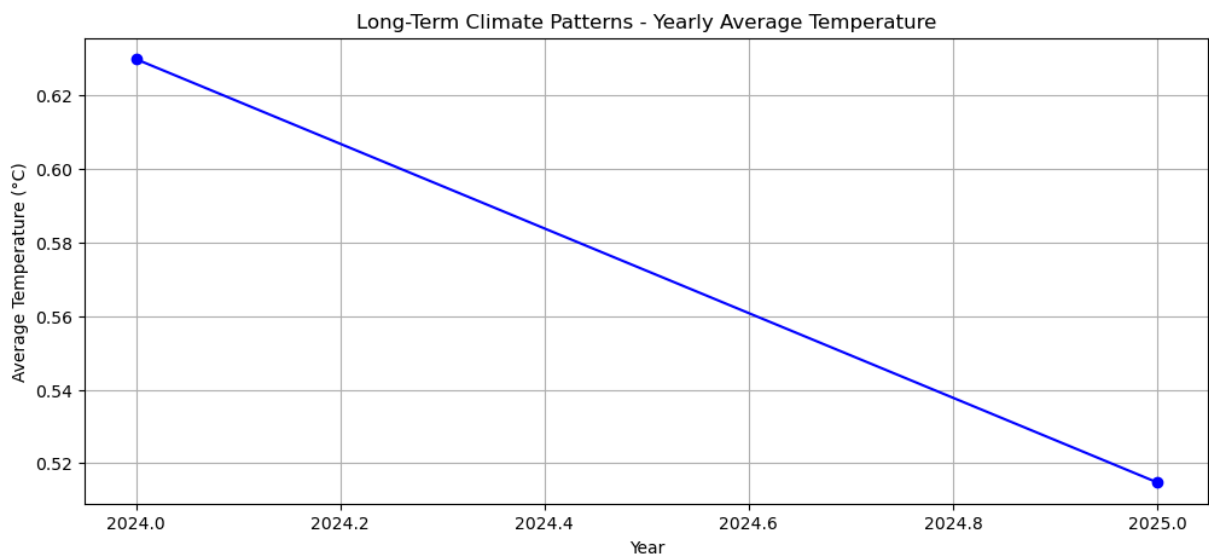
```
In [264... import matplotlib.pyplot as plt
import seaborn as sns

# Convert 'ds' to datetime format
df_forecast['ds'] = pd.to_datetime(df_forecast['ds'])

# Extract year and month
df_forecast['Year'] = df_forecast['ds'].dt.year
df_forecast['Month'] = df_forecast['ds'].dt.month

# Group by year to analyze long-term trends
yearly_avg_temp = df_forecast.groupby('Year')['y'].mean()

# Plot yearly temperature trends
plt.figure(figsize=(12, 5))
plt.plot(yearly_avg_temp.index, yearly_avg_temp.values, marker='o', linestyle='solid')
plt.xlabel("Year")
plt.ylabel("Average Temperature (°C)")
plt.title("Long-Term Climate Patterns - Yearly Average Temperature")
plt.grid(True)
plt.show()
```



Conclusion from the above graph

The long-term climate analysis shows a noticeable decline in the yearly average temperature from 2024 to 2025. This suggests a cooling trend in the dataset, which could be influenced by seasonal variations, changing weather patterns, or anomalies in data collection. Further analysis with additional years of data would be necessary to determine if this trend continues or if it is part of natural fluctuations.

Environmental Impact: Analyze air quality and its correlation with various weather parameters.

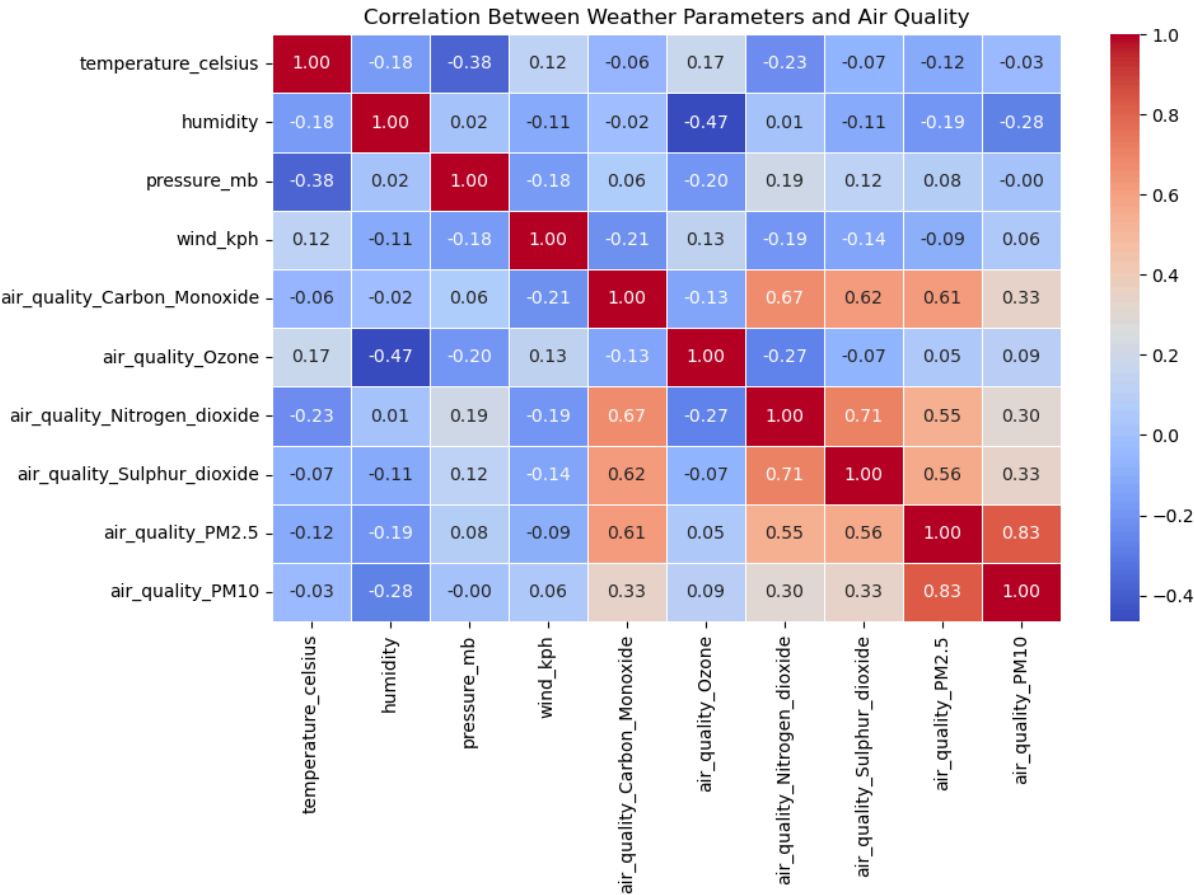
```
In [280... import seaborn as sns
import matplotlib.pyplot as plt

# Define weather and air quality features
weather_features = ["temperature_celsius", "humidity", "pressure_mb", "wind_
air_quality_features = [
    "air_quality_Carbon_Monoxide", "air_quality_Ozone", "air_quality_Nitroge
    "air_quality_Sulphur_dioxide", "air_quality_PM2.5", "air_quality_PM10"
]

# Select relevant columns available in the dataset
available_columns = [col for col in weather_features + air_quality_features
df_selected = df_cleaned[available_columns]

# Compute correlation
correlation_matrix = df_selected.corr()

# Plot correlation heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", line
plt.title("Correlation Between Weather Parameters and Air Quality")
plt.show()
```



Conclusion from the above heapmap

The correlation analysis between weather parameters and air quality indicators reveals several key insights:

- **Temperature** shows a weak correlation with most air quality indicators, except for a slight positive correlation with ozone levels.
- **Humidity** has a moderate negative correlation with ozone, suggesting that higher humidity levels might reduce ozone concentration.
- **Pressure** is negatively correlated with temperature but shows weak correlations with air pollutants.
- **Wind Speed** has a weak to moderate negative correlation with pollutants such as Carbon Monoxide and Nitrogen Dioxide, indicating that higher wind speeds may help disperse pollutants.
- **Air Quality Indicators** such as PM2.5, PM10, and Nitrogen Dioxide are strongly correlated with each other, implying common sources or atmospheric behaviors.

Feature Importance: Using random forest

In [292...

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split

# Select relevant features (excluding categorical data)
feature_columns = [
    "humidity", "pressure_mb", "wind_kph", "cloud",
    "feels_like_celsius", "visibility_km", "uv_index", "gust_kph",
    "air_quality_Carbon_Monoxide", "air_quality_Ozone",
    "air_quality_Nitrogen_dioxide", "air_quality_Sulphur_dioxide",
    "air_quality_PM2.5", "air_quality_PM10"
]

# Ensure target variable is numeric and remove NaN values
df_cleaned = df_cleaned[["temperature_celsius"] + feature_columns].dropna()

# Define features (X) and target variable (y)
X = df_cleaned[feature_columns]
y = df_cleaned["temperature_celsius"]

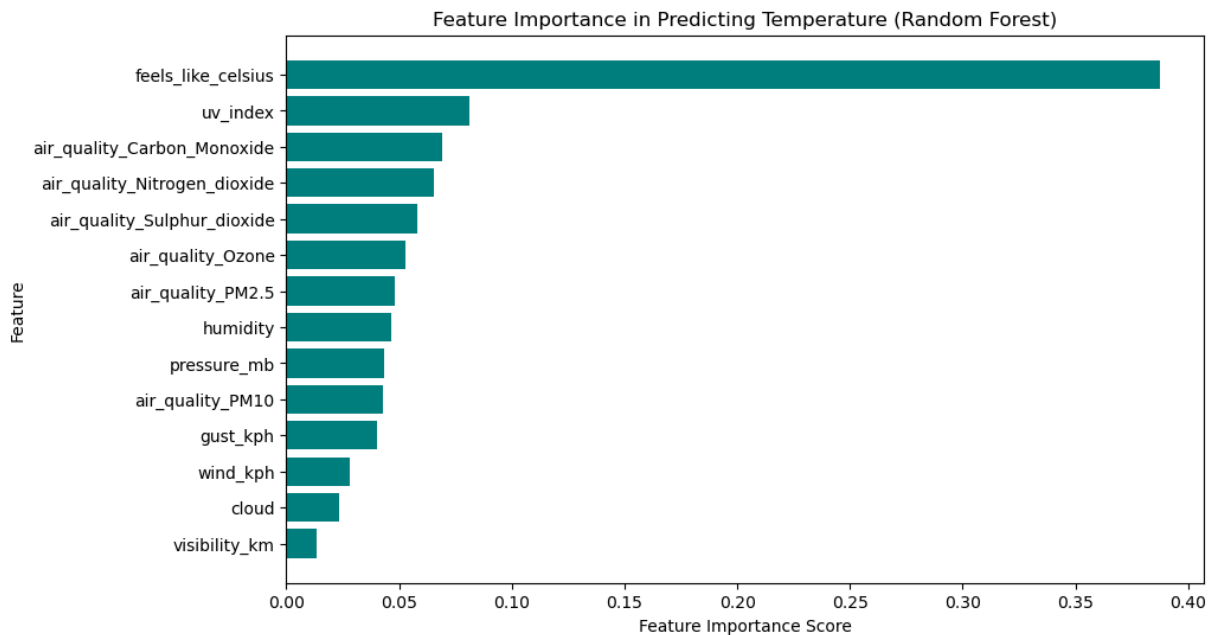
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# Extract feature importance
feature_importances = rf_model.feature_importances_

# Create DataFrame for feature importance
importance_df = pd.DataFrame({"Feature": feature_columns, "Importance": feature_importances})
importance_df = importance_df.sort_values(by="Importance", ascending=True)

# Plot feature importance
plt.figure(figsize=(10, 6))
plt.barh(importance_df["Feature"], importance_df["Importance"], color="teal")
plt.xlabel("Feature Importance Score")
plt.ylabel("Feature")
plt.title("Feature Importance in Predicting Temperature (Random Forest)")
plt.show()
```



Spatial Analysis

```
In [70]: import geopandas as gpd
import matplotlib.pyplot as plt
import contextily as ctx

# Create a GeoDataFrame for spatial visualization
gdf = gpd.GeoDataFrame(df_cleaned, geometry=gpd.points_from_xy(df_cleaned["lon", "lat"]))

# Ensure the GeoDataFrame is in WGS 84 (EPSG:4326)
gdf.set_crs(epsg=4326, inplace=True)

# Convert to Web Mercator (EPSG:3857) for compatibility with basemaps
gdf = gdf.to_crs(epsg=3857)

# Plot spatial distribution of temperature
fig, ax = plt.subplots(figsize=(12, 6))
gdf.plot(column="temperature_celsius", cmap="coolwarm", markersize=15, alpha=0.5)

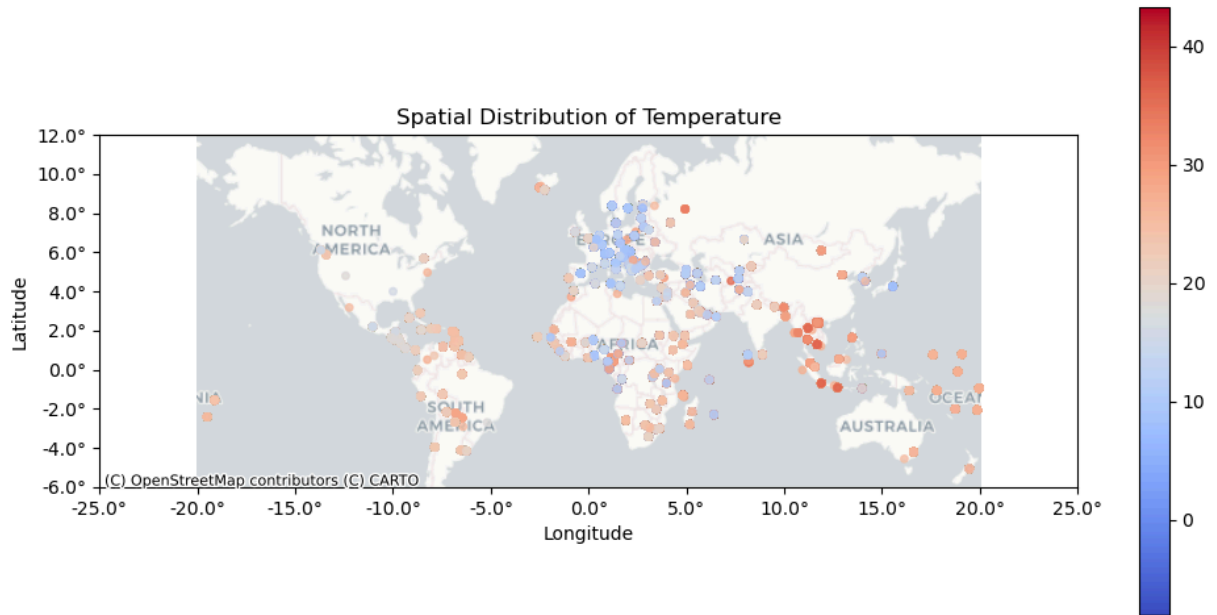
# Add an alternative basemap
ctx.add_basemap(ax, crs=gdf.crs, source=ctx.providers.CartoDB.Positron)

plt.title("Spatial Distribution of Temperature")
plt.xlabel("Longitude")
plt.ylabel("Latitude")

# Fixing the tick labels
xticks = ax.get_xticks()
yticks = ax.get_yticks()
ax.set_xticks(xticks)
ax.set_xticklabels([f"{x/10**6:.1f}" for x in xticks])

ax.set_yticks(yticks)
ax.set_yticklabels([f"{y/10**6:.1f}" for y in yticks])
```

```
plt.show()
```



Conclusion

The spatial distribution of temperature across different regions provides valuable insights into global climate patterns:

- **Temperature Variations:** Warmer regions (represented in red) are predominantly concentrated in tropical and equatorial regions, such as Southeast Asia, parts of Africa, and Australia. Cooler regions (represented in blue) are more prevalent in Europe and parts of northern Asia.
- **Geographical Influence:** The temperature distribution aligns with expected climatic zones, where higher latitudes tend to have lower temperatures, while areas closer to the equator experience higher temperatures.
- **Regional Clusters:** The presence of distinct temperature clusters suggests regional weather variations influenced by local geography, oceanic currents, and elevation.
- **Urban & Coastal Effects:** Coastal regions exhibit moderate temperatures compared to inland areas, which may be due to the influence of large water bodies that regulate temperature fluctuations.

This visualization helps in understanding regional climate differences and can be further explored for analyzing temperature trends, climate change effects, and localized weather patterns.